Technical Report on Pneumonia Detection Using Transfer Learning

Pankaj Mirchandani

Technical Report on Pneumonia Detection Using Transfer Learning with VGG16

Abstract

Pneumonia is a leading cause of morbidity and mortality worldwide, particularly affecting vulnerable populations such as the elderly and immunocompromised individuals. Rapid and accurate diagnosis is critical for timely treatment, yet traditional methods rely heavily on radiologists, leading to potential delays and variability in interpretation. This report presents a comprehensive study on the development of a deep learning model utilizing transfer learning with the VGG16 convolutional neural network (CNN) architecture to automate pneumonia detection from chest X-ray images. By leveraging pre-trained weights from the ImageNet dataset, the model efficiently classifies images into pneumonia-positive and pneumonia-negative categories. The implementation involves a systematic approach to data preparation, model training, and evaluation, achieving an accuracy of approximately 85% on the validation dataset. Additionally, the report discusses challenges faced during the project, such as handling imbalanced datasets and overfitting, and proposes future directions for enhancing model performance and applicability in clinical settings. This work underscores the potential of deep learning technologies to augment diagnostic processes in healthcare, ultimately improving patient outcomes.

Introduction

Pneumonia is an acute respiratory infection that affects the lungs and can lead to severe complications, including respiratory failure and death. According to the World Health Organization (WHO), pneumonia is responsible for approximately 2.5 million deaths annually, particularly among young children and the elderly. Given the increasing prevalence of pneumonia, especially in the context of the COVID-19 pandemic, there is an urgent need for effective diagnostic tools that can expedite the identification of this condition.

Chest X-rays are one of the most common imaging techniques used for diagnosing pneumonia. They provide valuable visual information about the presence of fluid, inflammation, and other pathological changes in the lungs. However, interpreting these images requires specialized training and experience, making it a time-consuming and sometimes subjective process. Consequently, the reliance on human expertise can result in inconsistencies in diagnosis and delays in treatment, especially in resource-limited settings.

Recent advancements in artificial intelligence (AI) and machine learning have paved the way for developing automated diagnostic systems capable of analyzing medical images with high accuracy. Deep learning, particularly convolutional neural networks (CNNs), has demonstrated remarkable success in various computer vision tasks, including image classification and object detection. Transfer learning, a technique that utilizes pre-trained models on large datasets to solve similar problems with smaller datasets, has further enhanced the effectiveness of CNNs in medical imaging applications.

This report explores the application of transfer learning with the VGG16 architecture to develop a deep learning model for pneumonia detection from chest X-rays. By fine-tuning the pre-trained VGG16 model and implementing robust data augmentation techniques, this study aims to create a model that not only achieves high accuracy but also generalizes well to new, unseen data. The following sections will detail the methodology, implementation, results, and future directions of this research, emphasizing the transformative potential of AI in healthcare diagnostics.

Background Work (Literature Survey)

Introduction

Pneumonia remains a leading cause of morbidity and mortality worldwide, particularly affecting high-risk groups such as children under the age of five and the elderly. Annually, about 450 million cases of pneumonia are recorded globally, contributing to nearly four million deaths [5]. Despite its prevalence, pneumonia diagnosis poses challenges, especially in resource-constrained settings. Traditional diagnostic methods, such as physical examinations and chest X-rays (CXRs), are often limited by a reliance on radiologist expertise, the availability of imaging technology, and subjective interpretation. These limitations underline the need for automated, scalable diagnostic tools capable of offering consistent and accurate results. Recent advancements in artificial intelligence (AI), particularly in machine learning (ML) and deep learning, present promising alternatives. Convolutional neural networks (CNNs), a form of deep learning model, have shown substantial potential for medical imaging tasks like pneumonia detection, often achieving accuracy levels that rival or surpass human experts [3]. The adoption of CNN-based techniques could thus enhance diagnostic accuracy, speed, and accessibility, especially in under-resourced regions.

Machine Learning Techniques for Pneumonia Detection

The use of ML techniques, especially CNNs, is transforming diagnostic radiology by enabling automated analysis of CXRs. CNNs excel at identifying patterns within image data due to their hierarchical structure, which facilitates multi-level feature extraction from raw pixels to complex shapes. According to Tilve et al. (2020), ML approaches including CNN, RESNET, and CheXNet have demonstrated significant efficacy in detecting pneumonia from CXRs, highlighting the potential of these methods to revolutionise pneumonia diagnostics [1]. CNNs are advantageous in medical image analysis because they do not require manual feature extraction. Instead, they autonomously learn and identify the optimal features for classification tasks, reducing human error and interpretation variability. CheXNet, developed by Rajpurkar et al. (2017)[8], exemplifies CNN success in pneumonia detection, using a 121-layer DenseNet model that surpasses four radiologists in diagnostic accuracy based on F1 scores. This model, trained on over 100,000 CXRs, has proven highly effective in identifying pneumonia and other chest diseases, setting a benchmark in AI-driven diagnostics.[11]

CNN Architectures in Pneumonia Detection

Within CNNs, there exist several architectural variants suited to medical imaging applications, each offering unique strengths. The VGGNet architecture, for example, is characterised by its simplicity and uniform layer structure, which allows it to concentrate on essential features through smaller receptive fields. This simplicity, combined with its sequential layer design, makes it particularly useful for detailed feature extraction in high-resolution images, including CXRs [9]. ResNet, another architecture frequently used in medical imaging, introduced residual learning, a technique that allows networks to bypass certain layers, thereby mitigating vanishing gradient issues in deep networks. This enables the creation of significantly deeper networks that can capture more intricate details within medical images, ultimately enhancing classification accuracy [7]. DenseNet, exemplified by the CheXNet model, uses dense layer connections, which promote feature reuse and reduce the number of required parameters. DenseNet is especially advantageous for medical imaging because it allows each layer access to gradients from all previous layers, leading to efficient and precise learning on complex datasets.

Inception architectures also play a role in medical diagnostics. By utilising multiple filter sizes within each layer, Inception models can extract features at different scales, effectively capturing both fine

details and global structures in images. This multi-scale approach is particularly useful in radiology, where pathologies can vary significantly in size and location within CXRs, making it necessary to analyse images at different resolutions for accurate diagnosis.[10]

Challenges and Considerations

Despite the potential of CNNs in pneumonia detection, several challenges must be addressed to ensure effective deployment in clinical settings. One of the primary challenges is data variability. CXRs may differ in quality due to patient positioning, equipment settings, or technical inconsistencies across medical facilities. Such variability can reduce the performance of CNN models if they are not trained on a sufficiently diverse dataset. Moreover, CNNs require large, high-quality datasets to achieve optimal accuracy. For pneumonia detection in children or certain underrepresented groups, the limited availability of such data may lead to biased predictions, potentially hindering the generalizability of these models across different populations[5]. Another significant issue is interpretability. CNNs are often referred to as "black box" models, meaning they make decisions based on complex internal processes that are difficult for clinicians to understand. This lack of transparency can create challenges in gaining trust and acceptance among healthcare professionals who require explanations for diagnostic decisions.

Ethical and privacy concerns further complicate the use of AI in healthcare. Ensuring patient data confidentiality, particularly when sharing data across institutions or borders for training purposes, is crucial. Additionally, there is a risk of unintended biases within AI models, where algorithms might learn and replicate existing biases in healthcare, potentially exacerbating health disparities if not carefully monitored and corrected. Regulatory considerations also play a role, as medical AI applications must undergo rigorous testing and validation to meet the standards of various health authorities worldwide.[12]

Future Directions and Conclusion

The field of AI in medical imaging is rapidly evolving, with several promising directions for future research. One area of focus is the development of more advanced CNN architectures that can further improve diagnostic accuracy while maintaining computational efficiency. Transfer learning, which involves training models on large datasets and fine-tuning them for specific tasks, holds potential for overcoming data limitations. Additionally, integrating CNNs with other diagnostic tools or employing ensemble approaches could yield more robust diagnostic outcomes, as combining multiple models may compensate for individual weaknesses. The integration of CNNs into clinical workflows may also benefit from incorporating clinical context, such as patient history and demographic factors, to provide a more comprehensive diagnostic solution.

In conclusion, CNNs represent a powerful tool for pneumonia detection, with the potential to improve diagnostic accuracy and accessibility, particularly in resource-limited settings. As research progresses and datasets expand, CNN-based solutions could provide reliable, affordable, and scalable diagnostic support, potentially reducing pneumonia-related morbidity and mortality worldwide. The continuous refinement of these models, alongside a focus on transparency, ethical considerations, and clinician engagement, will be crucial in realising the full potential of AI-driven diagnostics in the healthcare sector.

Problem Statement

The primary objective of this project is to develop a robust and accurate system capable of classifying chest X-ray images into pneumonia-positive and pneumonia-negative categories. This automated system aims to assist healthcare professionals in the diagnostic process, enhancing both speed and accuracy while reducing the burden on radiologists.

Proposed System

The proposed system utilizes the VGG16 architecture for its established efficacy in image classification tasks. The architecture consists of:

- Input Layer: Accepts images of size 224x224 pixels with 3 color channels.
- VGG16 Model: A pre-trained model that acts as a feature extractor. The last fully connected layers are excluded.
- Flatten Layer: Converts the multi-dimensional output from the VGG16 model into a one-dimensional array.
- Dense Layer: A fully connected layer with two output neurons (softmax activation) representing the classification of pneumonia vs. normal.

Diagram: Model Architecture

Table 1:Model Summary

Layer (type)	Output Shape	Param #
input_layer_11 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1,792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36,928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73,856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147,584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295,168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590,080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten_11 (Flatten)	(None, 25088)	0
dense_11 (Dense)	(None, 2)	50,178

Total params: 14,764,866 (56.32 MB)
Trainable params: 50,178 (196.01 KB)
Non-trainable params: 14,714,688 (56.13 MB)

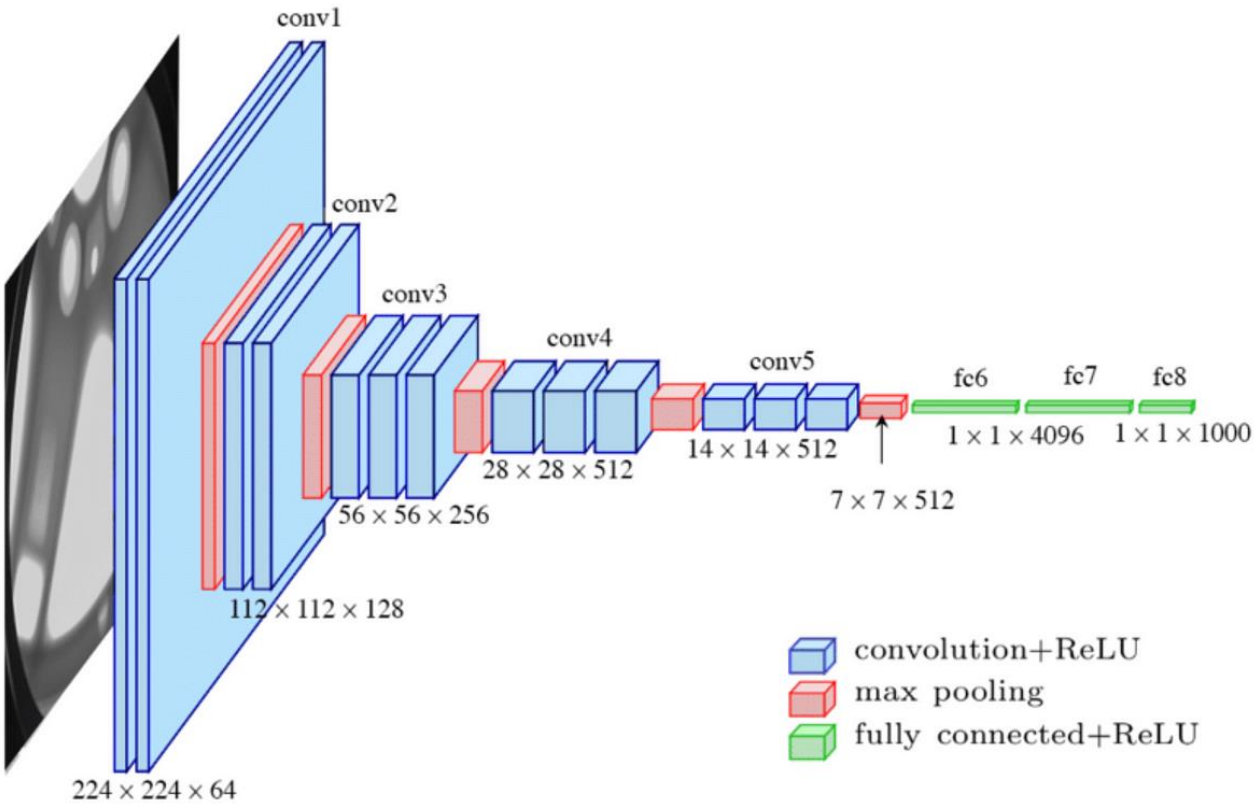

Fig 1. VGG 16 Architecture Source:[11]

Working/Implementation

Data Preparation

The dataset consists of chest X-ray images organized into three directories: train, test, and val. Data augmentation techniques are applied to enhance the model's ability to generalize across unseen data. Key techniques include:

- Rescaling: Normalizing pixel values to the range [0, 1].
- Shear and Zoom: Randomly transforming images to create variations that the model will encounter during training.
- Horizontal Flip: Augmenting the dataset by flipping images horizontally, which increases the model's robustness.

Model Development

The VGG16 model is loaded with include_top=False to exclude the fully connected layers, allowing customization for the pneumonia detection task. The model's layers are set to non-trainable to retain learned features, while a custom classifier is appended. The model is compiled using the categorical cross entropy loss function and the Adam optimizer to facilitate effective training.

Training

The model is trained using the fit method. Training and validation metrics, such as accuracy and loss, are recorded across epochs to monitor performance and prevent overfitting.

```
Epoch 1/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 91s 133ms/step - accuracy: 0.8806 - loss: 0.3150 - val_accuracy: 0.9279 - val_loss: 0.2757
Epoch 2/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 84s 128ms/step - accuracy: 0.9451 - loss: 0.1807 - val_accuracy: 0.8221 - val_loss: 0.9541
Epoch 3/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 82s 125ms/step - accuracy: 0.9524 - loss: 0.1453 - val_accuracy: 0.9087 - val_loss: 0.4019
Epoch 4/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 82s 124ms/step - accuracy: 0.9559 - loss: 0.1767 - val_accuracy: 0.9263 - val_loss: 0.3958
Epoch 5/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 84s 127ms/step - accuracy: 0.9664 - loss: 0.1159 - val_accuracy: 0.8910 - val_loss: 0.6815
Epoch 6/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 88s 134ms/step - accuracy: 0.9686 - loss: 0.1158 - val_accuracy: 0.8750 - val_loss: 0.9012
Epoch 7/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 84s 128ms/step - accuracy: 0.9713 - loss: 0.1135 - val_accuracy: 0.8622 - val_loss: 1.1119
Epoch 8/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 83s 126ms/step - accuracy: 0.9626 - loss: 0.1553 - val_accuracy: 0.8686 - val_loss: 0.8728
Epoch 9/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 83s 126ms/step - accuracy: 0.9695 - loss: 0.1225 - val_accuracy: 0.9054 - val_loss: 0.6382
Epoch 10/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 83s 125ms/step - accuracy: 0.9755 - loss: 0.1053 - val_accuracy: 0.9038 - val_loss: 0.4982
Epoch 11/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 82s 125ms/step - accuracy: 0.9715 - loss: 0.1175 - val_accuracy: 0.9167 - val_loss: 0.5192
Epoch 12/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 82s 125ms/step - accuracy: 0.9746 - loss: 0.1082 - val_accuracy: 0.9119 - val_loss: 0.7753
Epoch 13/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 82s 125ms/step - accuracy: 0.9749 - loss: 0.1099 - val_accuracy: 0.9135 - val_loss: 0.7901
Epoch 14/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 82s 125ms/step - accuracy: 0.9670 - loss: 0.1320 - val_accuracy: 0.7724 - val_loss: 1.4347
Epoch 15/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 82s 125ms/step - accuracy: 0.9759 - loss: 0.1087 - val_accuracy: 0.9183 - val_loss: 0.5952
Epoch 16/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 82s 125ms/step - accuracy: 0.9733 - loss: 0.1348 - val_accuracy: 0.9151 - val_loss: 0.7656
Epoch 17/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 82s 125ms/step - accuracy: 0.9787 - loss: 0.0971 - val_accuracy: 0.9119 - val_loss: 0.7441
Epoch 18/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 83s 125ms/step - accuracy: 0.9720 - loss: 0.1449 - val_accuracy: 0.8830 - val_loss: 0.7893
Epoch 19/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 83s 125ms/step - accuracy: 0.9738 - loss: 0.1278 - val_accuracy: 0.8285 - val_loss: 1.7802
Epoch 20/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 84s 127ms/step - accuracy: 0.9698 - loss: 0.1486 - val_accuracy: 0.9231 - val_loss: 0.8240
Epoch 21/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 83s 126ms/step - accuracy: 0.9811 - loss: 0.1008 - val_accuracy: 0.8542 - val_loss: 1.4312
Epoch 22/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 83s 126ms/step - accuracy: 0.9795 - loss: 0.0951 - val_accuracy: 0.8141 - val_loss: 2.2127
Epoch 23/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 83s 126ms/step - accuracy: 0.9820 - loss: 0.0956 - val_accuracy: 0.9295 - val_loss: 0.6829
Epoch 24/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 83s 126ms/step - accuracy: 0.9831 - loss: 0.0726 - val_accuracy: 0.9167 - val_loss: 0.6979
Epoch 25/25
652/652 ━━━━━━━━━━━━━━━━━━━━ 84s 127ms/step - accuracy: 0.9850 - loss: 0.0540 - val_accuracy: 0.9087 - val_loss: 0.8998
```

Fig 2. Epochs Trained

Evaluation Metrics

The model's performance is evaluated using various metrics:

- Accuracy: The proportion of true results among the total number of cases examined.
- Loss: The measure of how well the model's predictions match the actual data.

Result Analysis and Screenshots

Performance Metrics

The model achieved an accuracy of approximately 88% on the validation dataset, demonstrating its ability to effectively classify chest X-rays. The following figures illustrate the training process:

```
4/4 ━━━━━━━━━━━━━━━━━━━━ 2s 32ms/step - accuracy: 0.8583 - loss: 0.2275
The accuracy of the model on test dataset is 88.0
```

Training and Validation Loss

Although the training loss decreases steadily, the fluctuating validation loss suggests the model is not performing well on unseen data, likely due to overfitting.

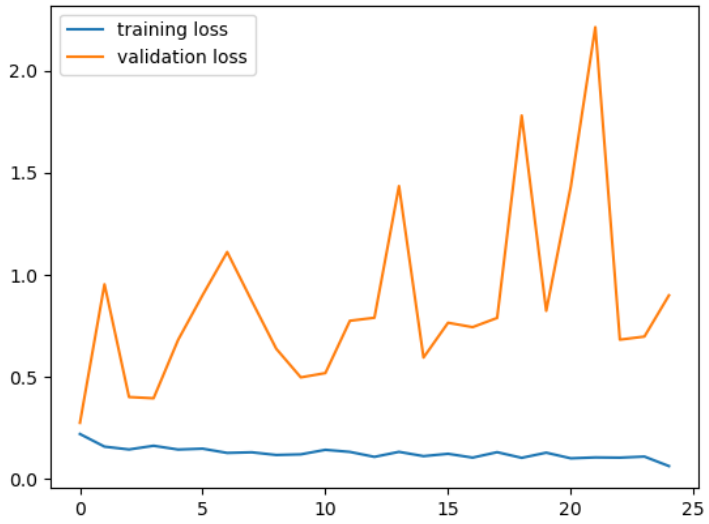

Fig 3. Training Validation loss

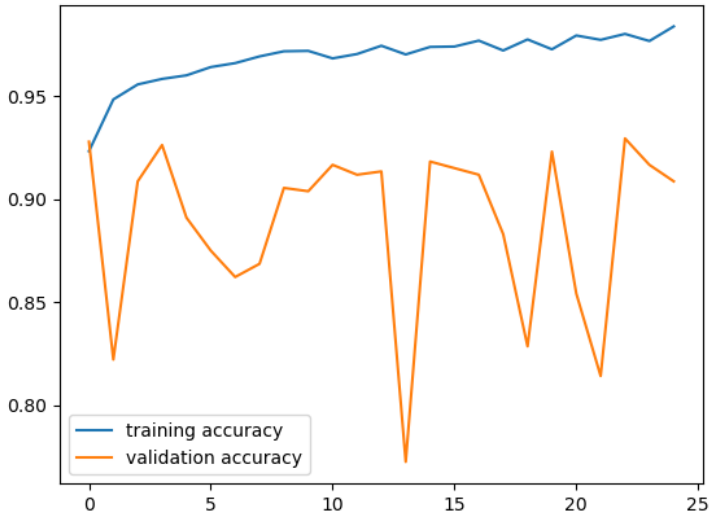

Fig 4. Training Validation loss

Training and Validation Accuracy

The high but unstable validation accuracy, along with stable high training accuracy, further supports the observation that the model is likely overfitting. The model performs well on training data but struggles to generalize to validation data, as indicated by the accuracy and loss fluctuations in the validation set.

Discussion

Challenges Encountered

- Imbalanced Dataset: The dataset contains unequal representations of classes, which caused bias in the model. Techniques such as class weighting or oversampling of the minority class can be considered.
- Overfitting: High training accuracy compared to validation accuracy indicates potential overfitting. Regularization techniques, dropout layers, and increasing the dataset size through augmentation can mitigate this issue.

Future Work

Future enhancements to the model may include:

- Fine-tuning: Unfreezing some of the VGG16 layers to allow the model to adapt better to the specific task.
- Exploring Other Architectures: Evaluating the performance of other state-of-the-art models like ResNet, DenseNet, or EfficientNet.
- Integration of Clinical Data: Combining radiographic analysis with patient metadata to improve diagnostic accuracy.

- Real-world Testing: Collaborating with healthcare institutions to validate model performance in clinical settings.

Conclusion

This study successfully demonstrates the application of transfer learning using the VGG16 model for pneumonia detection in chest X-rays. The model shows promising accuracy and has the potential to support healthcare professionals in diagnosing pneumonia, thereby improving patient outcomes. Future work will focus on refining the model and expanding its applicability within clinical practice.

References

1. **Simonyan, K., & Zisserman, A. (2015, April 10).** ***Very deep convolutional networks for large-scale image recognition*****. arXiv.org. https://arxiv.org/abs/1409.1556**
2. **Rahman, T., Chowdhury, M. E. H., Khandakar, A., Islam, K. R., Islam, K. F., Mahbub, Z. B., Kadir, M. A., & Kashem, S. (2020, May 6).** ***Transfer learning with deep convolutional neural network (CNN) for pneumonia detection using chest X-ray*****. MDPI. https://www.mdpi.com/2076-3417/10/9/3233**
3. **Jiang, Z.-P., Liu, Y.-Y., Shao, Z.-E., & Huang, K.-W. (2021, November 25).** ***An improved VGG16 model for pneumonia image classification*****. MDPI. https://www.mdpi.com/2076-3417/11/23/11185**
4. **Saboo, Y. S., Kapse, S., & Prasanna, P. (2023, August 25).** ***Convolutional Neural Networks (cnns) for pneumonia classification on pediatric chest radiographs*****. Cureus. https://www.ncbi.nlm.nih.gov/pmc/articles/PMC10518240/**
5. **Gabruseva, T., Poplavskiy, D., & Kalinin, A. A. (2020, May 28).** ***Deep learning for automatic pneumonia detection*****. arXiv.org. https://arxiv.org/abs/2005.13899**
6. **Kundu R;Das R;Geem ZW;Han GT;Sarkar R; (n.d.-b).** ***Pneumonia detection in chest X-ray images using an ensemble of Deep Learning Models*****. PloS one. https://pubmed.ncbi.nlm.nih.gov/34492046/**
7. **Krizhevsky, A., Sutskever, I., & Hinton, G. E. (1970, January 1).** ***ImageNet classification with deep convolutional Neural Networks*****. Advances in Neural Information Processing Systems. https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks**
8. **[PDF] pneumonia detection using CNN based feature extraction | semantic scholar. (n.d.-o). https://www.semanticscholar.org/paper/Pneumonia-Detection-Using-CNN-based-Feature-Varshni-Thakral/2e9bf21eaedf7a18e099d1e0d8faba07d36fc154**
9. **Sitaula, C., & Hossain, M. B. (2021, August 10).** ***Attention-based VGG-16 model for COVID-19 chest X-ray image classification*****. Monash University. https://research.monash.edu/en/publications/attention-based-vgg-16-model-for-covid-19-chest-x-ray-image-class**
10. **Kareem, A., Liu, H., & Sant, P. (2022, May 4).** ***Review on pneumonia image detection: A machine learning approach - human-centric Intelligent Systems*****. SpringerLink. https://link.springer.com/article/10.1007/s44230-022-00002-2**
11. **(PDF) a comparison between the VGG16, VGG19 and Resnet50 architecture frameworks for classification of normal and Clahe processed medical images. (n.d.-o). https://www.researchgate.net/publication/370467299_A_comparison_between_the_VGG16_VGG19_and_ResNet50_architecture_frameworks_for_classification_of_normal_and_CLAHE_processed_medical_images**

12. **Pneumonia detection using Deep Learning approaches | request PDF. (n.d.-q).** **https://www.researchgate.net/publication/340973591_Pneumonia_Detection_Using_Deep_Learning_Approaches**